

Introduction to Artificial Intelligence with Python

Learning


Supervised Learning

supervised learning

given a data set of input-output pairs, learn
a function to map inputs to outputs

classification

supervised learning task of learning a function mapping an input point to a discrete category



Date	Humidity (relative humidity)	Pressure (sea level, mb)	Rain

Date	Humidity (relative humidity)	Pressure (sea level, mb)	Rain
January 1	93%	999.7	Rain
January 2	49%	1015.5	No Rain
January 3	79%	1031.1	No Rain
January 4	65%	984.9	Rain
January 5	90%	975.2	Rain

f(humidity, pressure)

f(93, 999.7) = Rain


f(49, 1015.5) = No Rain


f(79, 1031.1) = No Rain


h(humidity, pressure)


pressure

humidity













nearest-neighbor classification


algorithm that, given an input, chooses the class of the nearest data point to that input
















k -nearest-neighbor classification

algorithm that, given an input, chooses the most common class out of the k nearest data points to that input





$x_1 = \text{Humidity}$

$x_2 = \text{Pressure}$

$h(x_1, x_2) = \begin{array}{l} \text{Rain if } w_0 + w_1x_1 + w_2x_2 \geq 0 \\ \text{No Rain otherwise} \end{array}$

Weight Vector \mathbf{w} : (w_0, w_1, w_2)

Input Vector \mathbf{x} : $(1, x_1, x_2)$

$$\mathbf{w} \cdot \mathbf{x}: w_0 + w_1x_1 + w_2x_2$$

$$h(x_1, x_2) = \begin{cases} 1 & \text{if } w_0 + w_1x_1 + w_2x_2 \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

Weight Vector \mathbf{w} : (w_0, w_1, w_2)

Input Vector \mathbf{x} : $(1, x_1, x_2)$

$$\mathbf{w} \cdot \mathbf{x}: w_0 + w_1x_1 + w_2x_2$$

$$h_{\mathbf{w}}(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x} \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

perceptron learning rule


Given data point (\mathbf{x}, y) , update each weight according to:


$$w_i = w_i + \alpha(y - h_{\mathbf{w}}(\mathbf{x})) \times x_i$$


perceptron learning rule

Given data point (\mathbf{x}, y) , update each weight according to:

$$w_i = w_i + \alpha(\text{actual value} - \text{estimate}) \times x_i$$








hard threshold

output

1

0

$w \cdot x$




soft threshold

output


1


0


$w \cdot x$



Support Vector Machines









maximum margin separator

boundary that maximizes the distance
between any of the data points





regression

supervised learning task of learning a function mapping an input point to a continuous value


f(advertising)

$$f(1200) = 5800$$

$$f(2800) = 13400$$

$$f(1800) = 8400$$

h(advertising)




Evaluating Hypotheses


loss function

function that expresses how poorly our hypothesis performs

0-1 loss function


$$L(\text{actual}, \text{predicted}) = \begin{cases} 0 & \text{if actual} = \text{predicted}, \\ 1 & \text{otherwise} \end{cases}$$






L_1 loss function

$$L(\text{actual}, \text{predicted}) = | \text{actual} - \text{predicted} |$$





L₂ loss function


$$L(\text{actual}, \text{predicted}) = (\text{actual} - \text{predicted})^2$$

overfitting

a model that fits too closely to a particular data set and therefore may fail to generalize to future data


pressure

humidity




pressure


humidity




pressure

humidity







$$\textit{cost}(\mathbf{h}) = \textit{loss}(\mathbf{h})$$

$$\textit{cost}(\mathbf{h}) = \textit{loss}(\mathbf{h}) + \textit{complexity}(\mathbf{h})$$

$$\textit{cost}(\mathbf{h}) = \textit{loss}(\mathbf{h}) + \lambda \textit{complexity}(\mathbf{h})$$

regularization

penalizing hypotheses that are more complex
to favor simpler, more general hypotheses

$$cost(h) = loss(h) + \lambda complexity(h)$$

holdout cross-validation

splitting data into a **training set** and a **test set**, such that learning happens on the training set and is evaluated on the test set

k -fold cross-validation

splitting data into k sets, and experimenting k times, using each set as a test set once, and using remaining data as training set

scikit-learn

Learning

- Supervised Learning
- Reinforcement Learning
- Unsupervised Learning

Supervised Learning

Introduction to Artificial Intelligence with Python