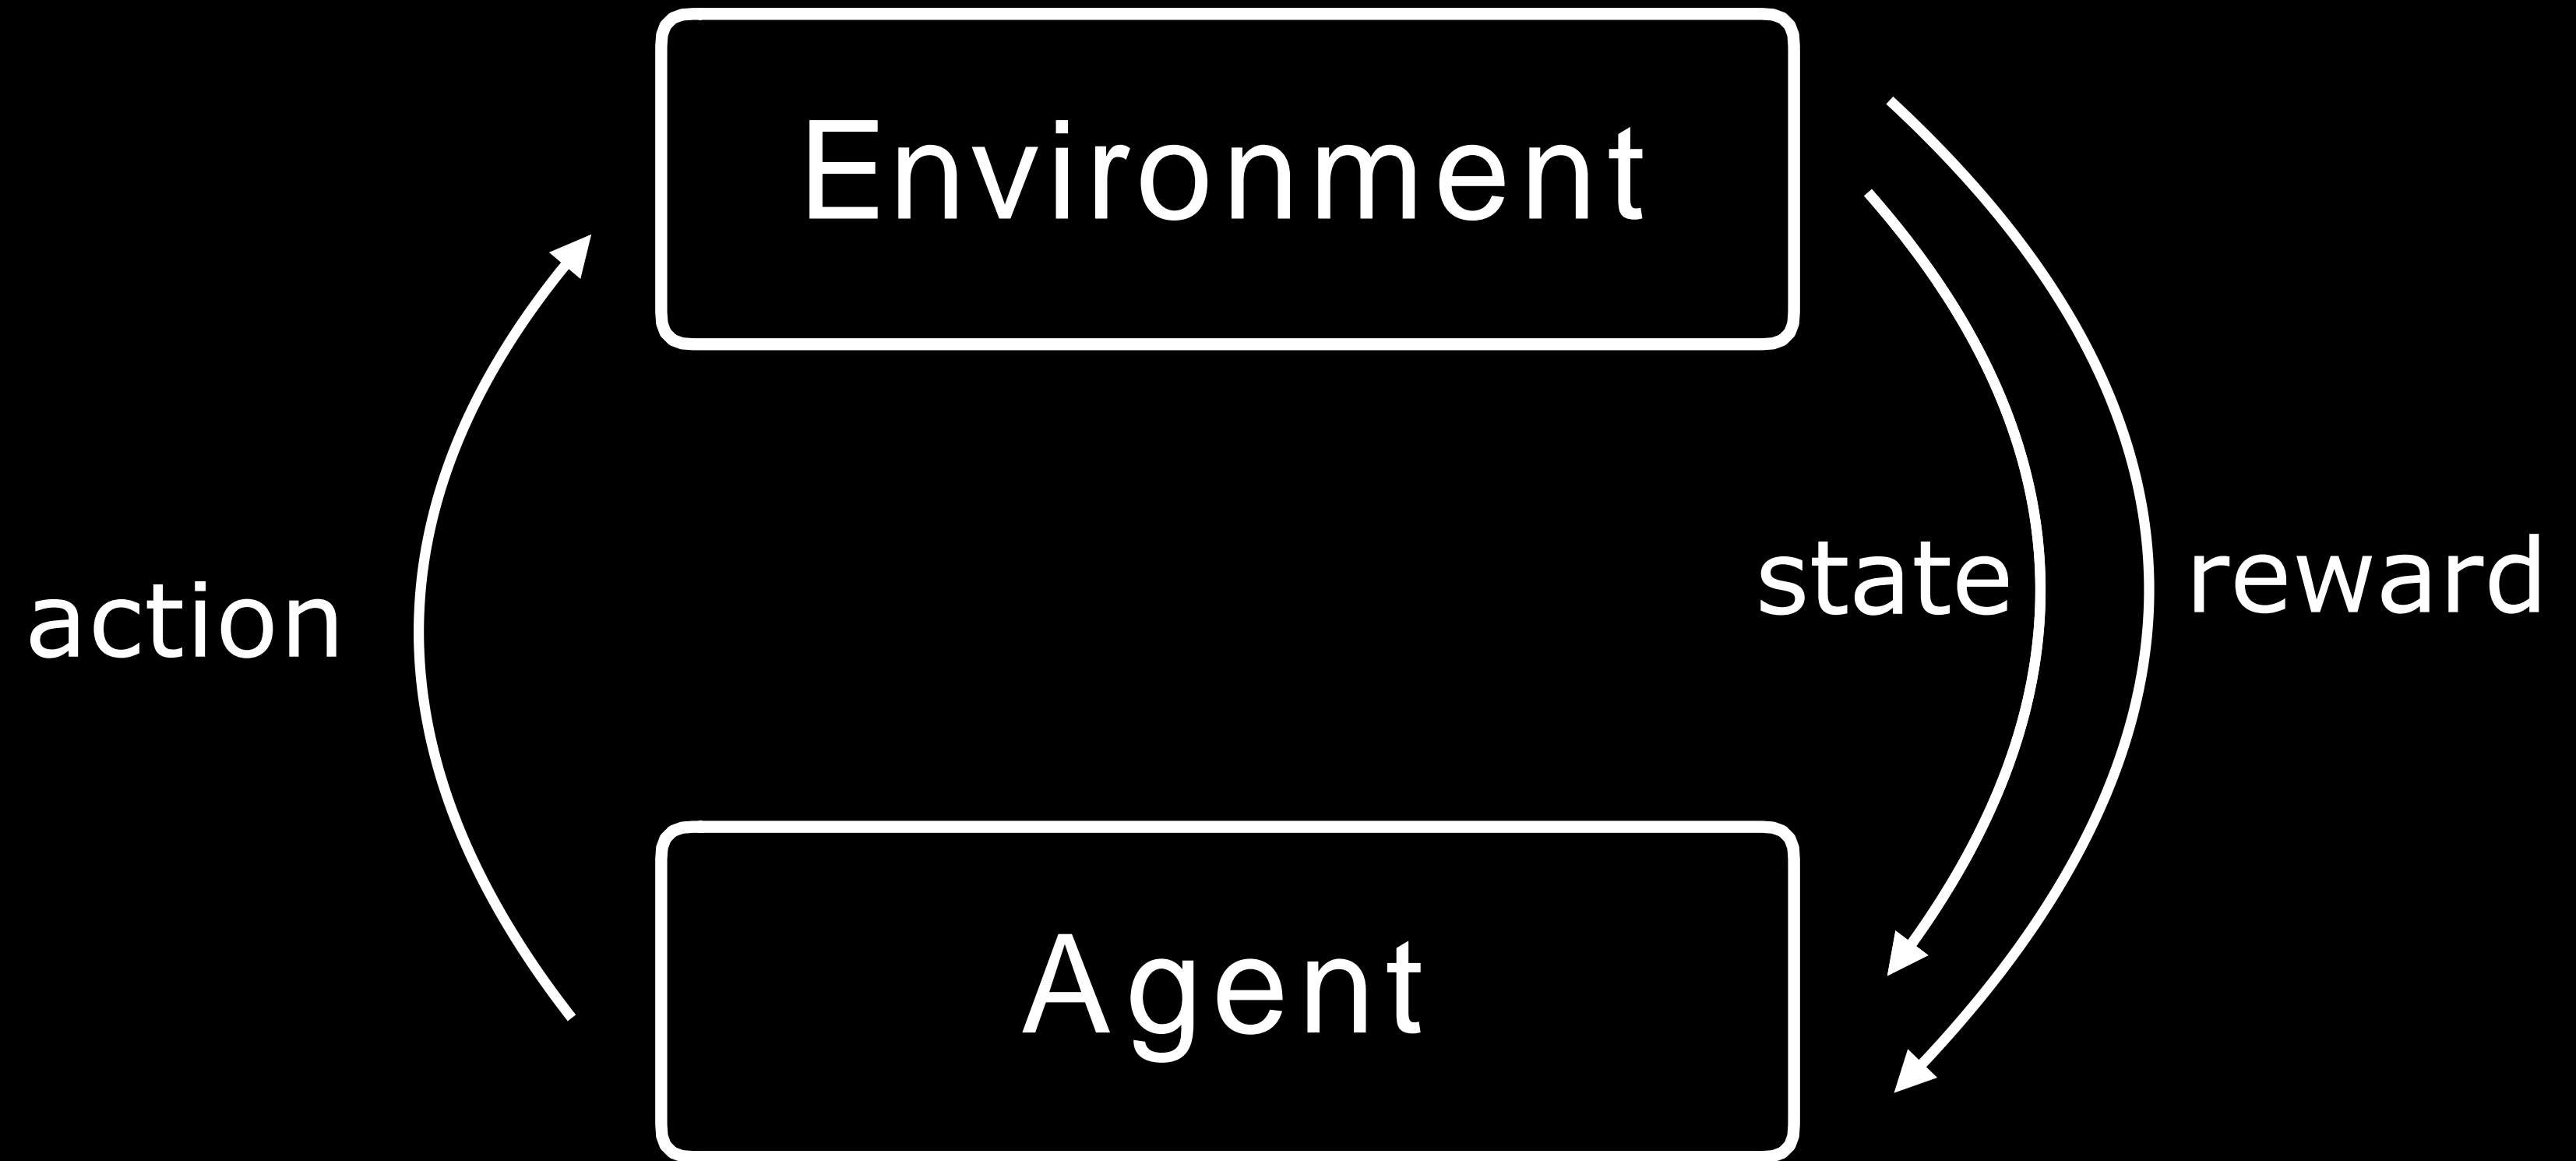


Introduction to Artificial Intelligence with Python

Reinforcement Learning

reinforcement learning

given a set of rewards or punishments, learn
what actions to take in the future




Markov Decision Process


model for decision-making, representing
states, actions, and their rewards


Markov Decision Process


model for decision-making, representing
states, actions, and their rewards

Markov Chain













Markov Decision Process


- Set of states S
- Set of actions $ACTIONS(s)$
- Transition model $P(s' | s, a)$
- Reward function $R(s, a, s')$




















Q-learning

method for learning a function $Q(s, a)$,
estimate of the value of performing action a
in state s

Q-learning Overview

- Start with $Q(s, a) = 0$ for all s, a
- When we taken an action and receive a reward:
 - Estimate the value of $Q(s, a)$ based on current reward and expected future rewards
 - Update $Q(s, a)$ to take into account old estimate as well as our new estimate

Q-learning

- Start with $Q(s, a) = 0$ for all s, a
- Every time we take an action a in state s and observe a reward r , we update:

$$Q(s, a) \leftarrow Q(s, a) + \alpha(\text{new value estimate} - \text{old value estimate})$$

Q-learning

- Start with $Q(s, a) = 0$ for all s, a
- Every time we take an action a in state s and observe a reward r , we update:

$$Q(s, a) \leftarrow Q(s, a) + \alpha(\text{new value estimate} - Q(s, a))$$

Q-learning

- Start with $Q(s, a) = 0$ for all s, a
- Every time we take an action a in state s and observe a reward r , we update:

$$Q(s, a) \leftarrow Q(s, a) + \alpha((r + \text{future reward estimate}) - Q(s, a))$$

Q-learning

- Start with $Q(s, a) = 0$ for all s, a
- Every time we take an action a in state s and observe a reward r , we update:

$$Q(s, a) \leftarrow Q(s, a) + \alpha((r + \max_{a'} Q(s', a')) - Q(s, a))$$


Q-learning


- Start with $Q(s, a) = 0$ for all s, a
- Every time we take an action a in state s and observe a reward r , we update:

$$Q(s, a) \leftarrow Q(s, a) + \alpha((r + \gamma \max_{a'} Q(s', a')) - Q(s, a))$$

Greedy Decision-Making

- When in state s , choose action a with highest $Q(s, a)$







Explore vs. Exploit


ϵ -greedy






- Set ϵ equal to how often we want to move randomly.
- With probability $1 - \epsilon$, choose estimated best move.
- With probability ϵ , choose a random move.




Nim


















function approximation

approximating $Q(s, a)$, often by a function combining various features, rather than storing one value for every state-action pair

Reinforcement Learning

Introduction to Artificial Intelligence with Python